

Data Science Advanced

Lesson02–Regression, Logistic Regression using Gradient Descent

Objective

After completing this lesson you will be able to:



Linear Regression:

- Describe hypothesis for a linear regression
- Understand cost function as a measure to derive regression equation.
- Understand gradient descent algorithm and its working to minimize the cost function.

Logistic Regression:


- Describe hypothesis for a logistic regression
- Understand cost function as a measure to derive logistic regression equation.
- Understand gradient descent algorithm and its working to minimize the cost function for a logistic regression

K Means Clustering:

- Understand cost function as a measure to derive K Means Cluster
- Understand Optimization objective for K means cluster

Linear Regression

Linear Regression to predict housing prices for a given size.



- It is a supervised learning problem as the “right answer” for each example is given in the dataset.
- A regression based supervised learning to predict real valued output.

Linear Regression–Data Representation

Size in feet ² (x)	Price (\$) in 1000's (y)
2104	460
1416	232
1534	315
852	178
...	...

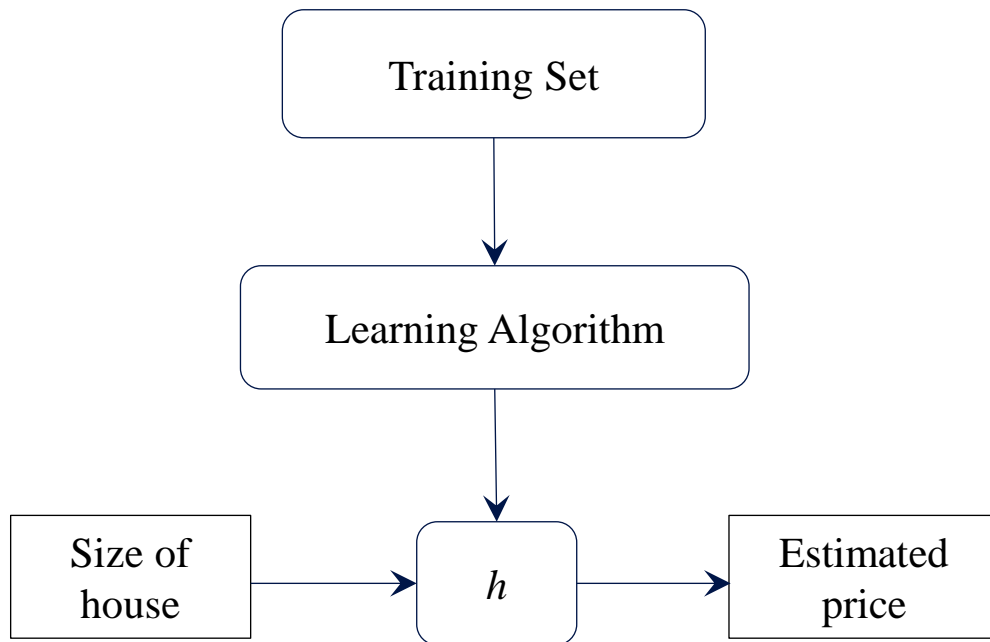
(x^i, y^i) – represents i^{th} training example

$x^1 = 2104$
 $y^1 = 460$
 $(x^1, y^1) = (2104, 460)$



Notions: **m** = Number of training examples
x's = “input” variable / features
y's = “output” variable / “target” variable

Regression–Hypothesis Formulation




- In case of Linear Regression, the hypothesis function is:

$$h_{\theta}(x) = \theta_0 + \theta_1 * x$$

- How to choose θ s?

Regression–Cost Function

Choose θ_0, θ_1 so that $h_{\theta}(x)$ is close to y for the training examples (x, y) .



The cost function:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^{(i)})^2$$

Goal: $\min_{\theta_0, \theta_1} (J(\theta_0, \theta_1))$




- Cost function $J(\theta_0, \theta_1)$ for regression is also called squared error function.
- The mean is halved ($1/(2*m)$) as a convenience for the computation of the gradient descent. The derivative term of the square function will cancel out the $1/2$ term.

Regression–Cost Function Intuition with one parameter


$$h_{\theta}(x)$$

(for fixed θ_1 , this is a function of x)



$$J(\theta_1)$$

(function of the parameter θ_1)




A simplified case where $\theta_0 = 0$. The cost function will be minimum for theta equal to zero. This will always be a convex shaped function.

Regression–Cost Function Intuition with two parameter


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$



(function of the parameters θ_0, θ_1)



Gradient Descent

Gradient Descent algorithm:

- Start with some θ_0, θ_1
- Keep changing θ_0, θ_1 to reduce $J(\theta_0, \theta_1)$ until a minimum is reached.



Gradient descent is a generic algorithm which can be used to minimize any type of cost function.

The initiation of θ_0, θ_1 can lead to a different local optima.

Gradient Descent

- The gradient descent algorithm is:

repeat until convergence{

$$\theta_j := \theta_j - \alpha * \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1)$$

}

- Simultaneous update of θ_0, θ_1 is needed:

$$\text{temp0} := \theta_0 - \alpha * \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\text{temp1} := \theta_1 - \alpha * \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\theta_1 := \text{temp1}$$



- α is the learning rate which decides how big or small the steps of descent will be.
- Near to the local minimum, the gradient descent will automatically take smaller steps.
- At the local optima, the θ_0, θ_1 does not change as derivative term will equal to zero.

Linear Regression–Gradient Descent

The gradient descent algorithm is:

repeat until convergence{
 $\theta_j := \theta_j - \alpha * \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ (for $j = 0$ and $j = 1$)
}

The cost function:

$$h_{\theta}(x) = \theta_0 + \theta_1 * x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^{(i)})^2$$

Goal: $\min_{\theta_0, \theta_1} (J(\theta_0, \theta_1))$


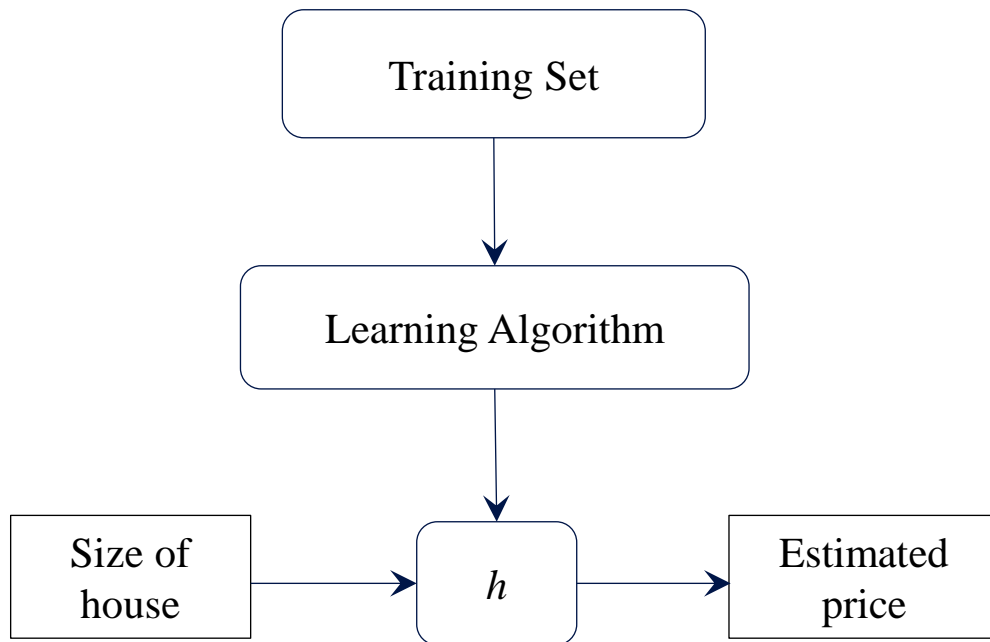
The derivate term for linear regression will be:

$$(\theta_0) J = 0: \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^{(i)})$$

$$(\theta_1) J = 1: \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^{(i)}) * x^i$$

Logistic Regression

Logistic Regression–Hypothesis Formulation



In case of Logistic Regression, the hypothesis function is:

$$h_{\theta}(x) = g(\theta_0 + \theta_1 * x)$$
$$g(z) = \frac{1}{1 + e^{-z}}$$

Where $z = \theta_0 + \theta_1 * x$


Logistic Regression–Decision Boundary

$$P(y = 1|x): h_{\theta}(x) = g(\theta_0 + \theta_1 * x)$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

Where $z = \theta_0 + \theta_1 * x$

- Predict $y = 1$ if $h_{\theta}(x) \geq 0.5$
- Predict $y = 0$ if $h_{\theta}(x) < 0.5$



$g(z) \geq 0.5$ whenever $z \geq 0$

Since

$$h_{\theta}(x) = g(\theta_0 + \theta_1 * x)$$

so


$$h_{\theta}(x) \geq 0.5 \text{ when } (\theta_0 + \theta_1 * x) \geq 0$$

Logistic Regression – Cost Function

- Choose θ_0, θ_1 so that $h_{\theta}(x)$ is close to y for the training examples (x, y) . Rewriting the linear regression cost function.

$$J(\theta) = 1/m \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^i) \quad \text{where } \text{Cost}(h_{\theta}(x^{(i)}), y^i) = 1/2 \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^i)^2$$



- Cost depicts the penalty learning algorithm has to pay when it outputs $h_{\theta}(x^{(i)})$ when actual label is y
- The above cost function cannot be used for logistic regression as it will be a non-convex function.



Logistic Regression – Cost Function

- The penalty, learning algorithm has to pay when it outputs $h_{\theta}(x^{(i)})$ when actual label is y is

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -(1 - \log(h_{\theta}(x))) & \text{if } y = 0 \end{cases}$$



Logistic Regression – Simplified Cost Function

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -(1 - \log(h_{\theta}(x))) & \text{if } y = 0 \end{cases}$$

The simplified cost function is

The cost function:

$$J(\theta_0, \theta_1) = -\frac{1}{m} \sum_{i=1}^m y^i * \log(h_{\theta}(x^i)) + (1 - y^{(i)}) * \log(1 - h_{\theta}(x^i))$$

Goal: $\min_{\theta_0, \theta_1} (J(\theta_0, \theta_1))$

Logistic Regression–Gradient Descent

- The gradient descent algorithm is:

repeat until convergence{

$$\theta_j := \theta_j - \alpha * \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1)$$

}

- Simultaneous update of θ_0, θ_1 is needed:

$$\text{temp0} := \theta_0 - \alpha * \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\text{temp1} := \theta_1 - \alpha * \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\theta_1 := \text{temp1}$$



- α is the learning rate which decides how big or small the steps of descent will be.
- Near to the local minimum, the gradient descent will automatically take smaller steps.
- At the local optima, the θ_0, θ_1 does not change as derivative term will equal to zero.

Logistic Regression–Gradient Descent

The gradient descent algorithm is:

```
repeat until convergence{  
     $\theta_j := \theta_j - \alpha * \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$   
    (for  $j = 0$  and  $j = 1$ )  
}
```

The cost function:

$$J(\theta_0, \theta_1) = -\frac{1}{m} \sum_{i=1}^m y^i * \log(h_{\theta}(x^i)) + (1 - y^{(i)}) * \log(1 - h_{\theta}(x^i))$$

Goal: $\min_{\theta_0, \theta_1} (J(\theta_0, \theta_1))$

The derivate term for logistic regression will be:

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^{(i)}) * x^{(i)}_j$$

K Means

K Means Algorithm

Randomly initialize K Cluster Centroids ($\mu_1, \mu_2, \mu_3 \dots \mu_k$)

Training set is $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots x^{(m)}\}$

Cluster assignment step

Repeat { $i = 1$ to m
 $c^{(i)}$ = index from 1 to K of cluster centroid closet to $x^{(i)}$
}

Move Centroid

Repeat { $k = 1$ to K
 μ_k = average of points assigned to cluster k
}

Cost function for K Means

$c^{(i)}$ = index of cluster (1, 2 , ... K) to which $x^{(i)}$ is currently assigned

μ_k = cluster centroid k where $k = \{1, 2, 3, \dots, K\}$

$\mu_{c^{(i)}}$ = cluster centroid of the cluster to which example $x^{(i)}$ has been assigned

Optimization objective

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_k) = 1/m \sum_1^m \|x^i - \mu_{c^{(i)}}\|^2$$

$$\underset{\substack{c^{(1)}, \dots, c^{(m)} \\ \mu_1, \dots, \mu_k}}{\text{minimize}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_k)$$

Random initialization of K clusters

Pick k random training examples and set μ_1, \dots, μ_k equal to these k training examples

```
for{ i = 1 to 100
```

```
    Randomly initialize K means
```

```
    Run K means to get cluster assignment and centroid movement
```


```
    Compute cost function
```

```
}
```

Pick the cluster that gives the lowest cost $J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_k)$

Optimal Number of Cluster

Elbow Method: Look for the number of cluster after which the delta change in cost function is not significant



Pick the cluster that gives the lowest cost $J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_k)$

End of Lesson02–Regression, Logistic Regression using Gradient Descent

