

german_credit_gradient_descent

June 9, 2018

Demonstration of Gradient Descent using R:

Kumar Rahul

In this exercise, we will use the German Credit Rating data to demonstrate Gradient Descent Algorithm. The Dataset can be downloaded from:

[https://archive.ics.uci.edu/ml/datasets/statlog+\(german+credit+data\)](https://archive.ics.uci.edu/ml/datasets/statlog+(german+credit+data))

1 Code starts here

We are going to use below mentioned libraries for demonstrating plot the cost function:

```
In [1]: library(ggplot2)
```

1.1 Data Import and Manipulation

1.1.1 1. Import a data set

```
In [2]: raw_data <- read.csv("/Users/Rahul/Documents/Datasets/german_credit_rating_data 2010.csv",
                             header = TRUE, sep = ",", na.strings = c("", " ", "NA"))

filter_data <- raw_data
filter_data <- raw_data[complete.cases(raw_data),]
filter_data <- raw_data[,c(-1)]
```

1.1.2 2. Create dummy variable for factors

```
In [3]: n <- ncol(filter_data)
        #for creating X0 variable
        processed_data <- cbind(rep(1, nrow(filter_data)))
        for (i in 1:n) {
          if (is.factor(filter_data[,i])) {
            #creating dummies for the factors and storing in temp
            temp <- model.matrix(~filter_data[,i])
            #removing the first column which is an intercept term in dummy coding
            temp <- subset(temp, select = -c(1) )
            #loop through to rename the column names appropriately
            for (c in 1:ncol(temp)) {
              str <- colnames(temp)[c]
```

```

        colnames(temp)[c] <- sub("filter_data\\[, i\\]", " ", colnames(temp)[c])
        colnames(temp)[c] <- paste(colnames(filter_data)[i], "- ", colnames(temp)[c])
    }
    #the factors after dummy variable creation is appended to the processed_data
    processed_data <- cbind(processed_data,temp)
}
#in case not a factor, store the variable in processed_data directly
else {
    processed_data <- cbind(processed_data,filter_data[,i])
    len <- ncol(processed_data)
    colnames(processed_data)[len] <- colnames(filter_data[i])

    #standardize the numeric variable for faster descent. (Feature scaling)
    processed_data[,len] <- scale(processed_data[,len])
}
}

```

1.1.3 3. Seperate x and y variable and converting to matrix form

```

In [4]: x <- processed_data
        y <- as.matrix(x[,ncol(x)])
        colnames(y)[1] <- colnames(x)[ncol(x)]
        x <- x[,c(-ncol(x))]
        colnames(x)[1] <- "X0"
        x <- as.matrix(x)

```

1.2 Create functions

1.2.1 1. Define Sigmoid function

```

In [5]: sigmoid <- function(z) {
        1/(1 + exp(-z))
    }

```

1.2.2 2. Define Hypothesis function

```

In [6]: hypothesisLr <- function(theta, x) {
        sigmoid(x %*% theta) #matrix multiplication (x is 1000*75 and theta is 75*1, result 75
    }

```

1.2.3 3. Define Cost function

```

In [7]: computeCost <- function(theta, x, y) {
        m <- length(y)
        s <- sapply(1:m, function(i)
        y[i]*log(hypothesisLr(theta,x[i,])) + (1 - y[i])*log(1 - hypothesisLr(theta,x[i,]))
        )
        j <- -1/m * sum(s)
    }

```

```

return(j)
}

```

1.2.4 4. Define Gradient Descent function

```

In [8]: gradAlgo <- function(theta, x, y) {
  m <- length(y)
  g <- 1/m * t(x) %*% (hypothesisLr(theta,x) - y)
  return(g)
}

```

1.3 Call functions

1.3.1 1. Initialize α, θ and compute cost

Run normal gradient descent to compute theta and cost function J for each iteration

```

In [9]: alpha <- 1 #learning rate 1
  iter <- 1000 #no of iterations for descent 5000
  #m <- length(y)
  theta <- rep(0, ncol(x))

  #cost at intial theta
  computeCost(theta, x, y)
  j <- rep(0,iter) #to maintain log of cost
  for (c in 1:iter) {
    theta <- theta - alpha * gradAlgo(theta,x,y)
    j[c] <- computeCost(theta,x,y)
  }

```

0.693147180559945


1.3.2 2. Plot Cost Value

Plot cost as a function of no of iterations to converge

```

In [10]: ggplot() +
  aes(x = 1:iter, y = j) +
  geom_point(colour = "red") +
  geom_path() + xlab("Iteration") +
  ylab("Cost J")

```



1.3.3 3. Build Classification matrix

```
In [11]: m <- nrow(y)
         actual_y <- y
         actual_y[actual_y == 0] <- "Bad"
         actual_y[actual_y == 1] <- "Good"
         predicted_y = rep("Bad", m);
         predicted_y[hypothesisLr(theta,x) >= 0.5] = "Good";
         addmargins(table(actual_y, predicted_y))
```

	Bad	Good	Sum
Bad	158	142	300
Good	73	627	700
Sum	231	769	1000

1.3.4 4. Score new cases

Scoring new profile pass all X's

```
In [12]: #new_score <- (sigmoid(c(1,<all X's>) %*% theta))
```

1.4 Compare the Gradient Descent with GLM

```
In [13]: lg_glm <- glm(filter_data$Credit.classification ~ .,data = filter_data, family = "binomial")
summary(lg_glm)
lg_predict_prob <- predict(lg_glm, filter_data[,1:20],"response")
m <- length(filter_data[,21])
actual_y <- filter_data[,21]

predicted_y = rep("Bad", m);
predicted_y[lg_predict_prob >= 0.5] = "Good";
addmargins(table(actual_y, predicted_y))
```

Call:

```
glm(formula = filter_data$Credit.classification ~ ., family = "binomial",
     data = filter_data)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.6116	-0.7095	0.3752	0.6994	2.3410

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	2.353e+00	1.330e+00	1.769	0.076938	.
CHK_ACCTless-200DM	3.749e-01	2.179e-01	1.720	0.085400	.
CHK_ACCTno-account	1.712e+00	2.322e-01	7.373	1.66e-13	***
CHK_ACCTover-200DM	9.657e-01	3.692e-01	2.616	0.008905	**
Duration	-2.786e-02	9.296e-03	-2.997	0.002724	**
History bank-paid-duly	-1.434e-01	5.489e-01	-0.261	0.793921	
History critical	1.436e+00	4.399e-01	3.264	0.001099	**
History delay	8.532e-01	4.717e-01	1.809	0.070470	.
History duly-till-now	5.861e-01	4.305e-01	1.362	0.173348	
Purpose.of.credit domestic-app	-2.173e-01	8.041e-01	-0.270	0.786976	
Purpose.of.credit education	-7.764e-01	4.660e-01	-1.666	0.095718	.
Purpose.of.credit furniture	5.152e-02	3.543e-01	0.145	0.884391	
Purpose.of.credit new-car	-7.401e-01	3.339e-01	-2.216	0.026668	*
Purpose.of.credit others	7.487e-01	7.998e-01	0.936	0.349202	
Purpose.of.credit radio-tv	1.515e-01	3.370e-01	0.450	0.653002	
Purpose.of.credit repairs	-5.237e-01	5.933e-01	-0.883	0.377428	
Purpose.of.credit retraining	1.319e+00	1.233e+00	1.070	0.284625	
Purpose.of.credit used-car	9.264e-01	4.409e-01	2.101	0.035645	*
Credit.Amount	-1.283e-04	4.444e-05	-2.887	0.003894	**
Balance.in.Savings.A.C less100DM	-3.761e-01	4.011e-01	-0.938	0.348476	

Balance.in.Savings.A.C less500DM	-1.833e-02	4.656e-01	-0.039	0.968595
Balance.in.Savings.A.C over1000DM	9.631e-01	6.425e-01	1.499	0.133868
Balance.in.Savings.A.C unknown	5.706e-01	4.492e-01	1.270	0.203940
Employment one-year	-1.159e-01	2.423e-01	-0.478	0.632415
Employment over-seven	9.379e-02	2.510e-01	0.374	0.708653
Employment seven-years	6.482e-01	2.684e-01	2.415	0.015734 *
Employment unemployed	-1.828e-01	4.105e-01	-0.445	0.656049
Install_rate	-3.301e-01	8.828e-02	-3.739	0.000185 ***
Marital.status male-divorced	-2.755e-01	3.865e-01	-0.713	0.476040
Marital.status married-male	9.162e-02	3.118e-01	0.294	0.768908
Marital.status single-male	5.406e-01	2.102e-01	2.572	0.010113 *
Co.applicant guarantor	1.415e+00	5.685e-01	2.488	0.012834 *
Co.applicant none	4.360e-01	4.101e-01	1.063	0.287700
Present.Resident	-4.776e-03	8.641e-02	-0.055	0.955920
Real.Estate car	8.690e-02	2.313e-01	0.376	0.707115
Real.Estate none	-4.490e-01	4.130e-01	-1.087	0.277005
Real.Estate real-estate	2.814e-01	2.534e-01	1.111	0.266630
Age	1.454e-02	9.222e-03	1.576	0.114982
Other.installment none	6.463e-01	2.391e-01	2.703	0.006871 **
Other.installment stores	1.232e-01	4.119e-01	0.299	0.764878
Residence own	-2.402e-01	4.503e-01	-0.534	0.593687
Residence rent	-6.839e-01	4.770e-01	-1.434	0.151657
Num_Credits	-2.721e-01	1.895e-01	-1.436	0.151109
Job skilled	-7.524e-02	2.845e-01	-0.264	0.791419
Job unemployed-non-resident	4.795e-01	6.623e-01	0.724	0.469086
Job unskilled-resident	-5.666e-02	3.501e-01	-0.162	0.871450
No..dependents	-2.647e-01	2.492e-01	-1.062	0.288249
Phone yes	-3.000e-01	2.013e-01	-1.491	0.136060
Foreign yes	-1.392e+00	6.258e-01	-2.225	0.026095 *

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1221.73 on 999 degrees of freedom
 Residual deviance: 895.82 on 951 degrees of freedom
 AIC: 993.82

Number of Fisher Scoring iterations: 5

	Bad	Good	Sum
bad.	160	140	300
good.	74	626	700
Sum	234	766	1000

End of Document
